

# *Comparative Analysis of Soft Actor-Critic and Proximal Policy Optimization Algorithms on the Walker2D Environment*

*Reinforcement Learning Final Project  
By Rufina George and Shrivatsa Mudligiri*


# Introduction

## Challenges in Robotics

- Traditional control methods in robotics struggle with complex and dynamic environments.
- These tasks demand precise and fine-grained control over robot actuators.
- Hard-coding behaviors for every task and environment becomes impractical.
- Dynamic and uncertain environments necessitate adaptive control strategies.

## The Solution: Reinforcement Learning

- Reinforcement Learning (RL) enables goal-oriented learning through interaction with the environment.
- Robots can learn optimal behaviors by exploring actions and learning from consequences.
- RL offers flexibility and adaptability, crucial for continuous control tasks.
- Continuous learning from interactions improves performance and robustness.
- Algorithms like Soft Actor-Critic (SAC) and Proximal Policy Optimization (PPO) are effective in continuous control.



# Introduction

## Benefits of Reinforcement Learning in Robotics


- Autonomous learning and adaptation without manual programming.
- Tackling complex and dynamic environments more effectively.

## Conclusion

- RL offers a powerful paradigm for achieving adaptive control in robotics.
- Advancements in RL have the potential to unlock new capabilities in robotic systems.

## Goal

- Evaluate and compare the performance of SAC and PPO on the Walker2D environment.
- Understand their strengths and weaknesses in the context of a complex control problem.
- Investigate factors such as training efficiency, stability, and robustness, which are important considerations in real-world applications of RL algorithms.



# Actor-critic methods

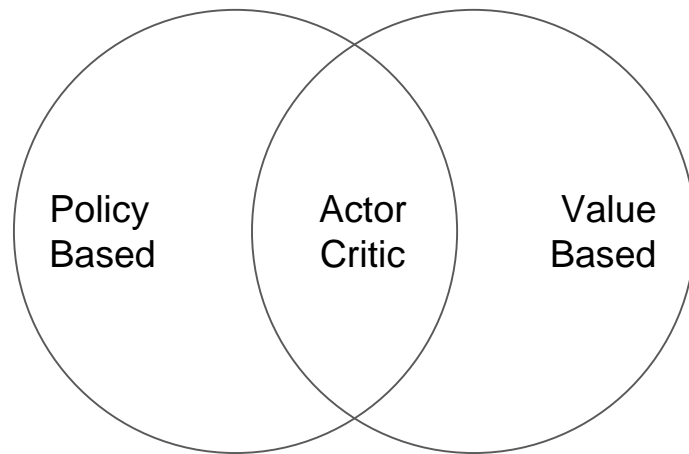
Actor-Critic methods combine value-based and policy-based approaches.

The “Critic” estimates the value function (action-value or state-value) and the “Actor” updates the policy distribution in the direction suggested by the Critic (such as with policy gradients).

Both the Critic and Actor functions are parameterized with neural networks.

## Advantages

- Efficient Exploration
- High Sample Efficiency
- Policy Improvement
- Flexible



# Soft Actor-Critic (SAC)

SAC combines (actor-critic + entropy regularization)

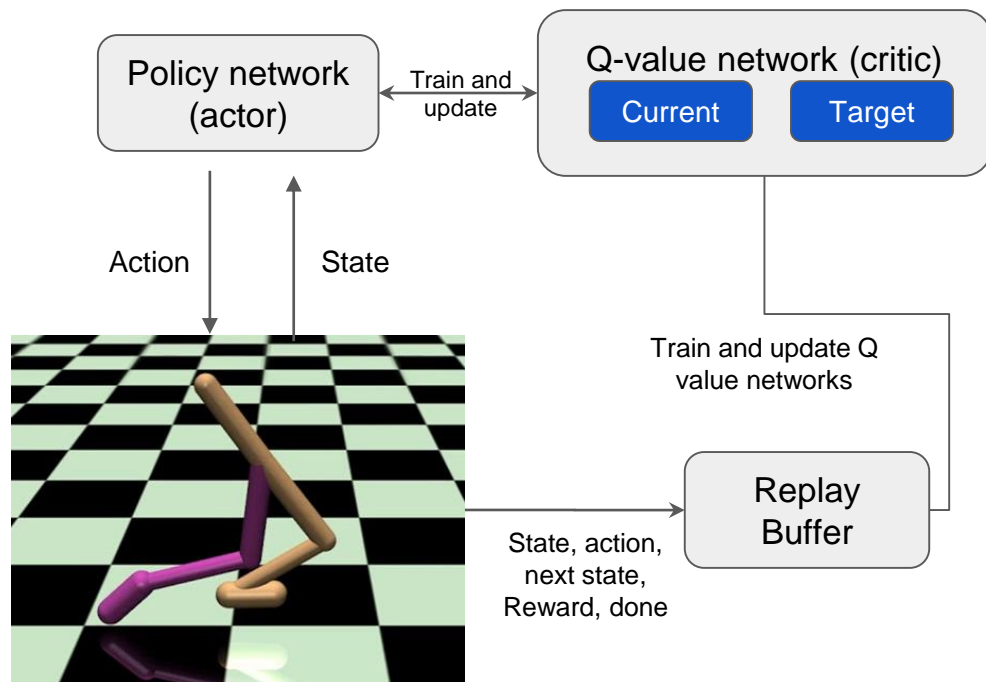
Key Features:

- Well-suited for tasks with continuous action spaces
- Directly optimizes the policy, learning the best actions to take in different situations.
- Estimates the value function to guide the learning process and evaluate the expected return or future rewards.
- Employs a soft value function update approach
- Soft Q-Learning and Soft Value Iteration use maximum entropy principles to encourage the learning of soft value functions.
- Policies achieve high rewards and capture the underlying uncertainty in the environment.

Applications:

- Successfully applied to a wide range of tasks, including robotic control, autonomous driving, and game playing.
- In robotics, SAC enables robots to perform precise and smooth movements, adapt to dynamic environments, and learn complex tasks.
- SAC has also shown promising results in complex game-playing domains, achieving impressive performance in games like Atari and Dota 2.

## Soft Actor-Critic: How it works?



Loss function for Q-network in SAC:

$$L(\phi_i, \mathcal{D}) = \mathbb{E}_{(s,a,r,s',d) \sim \mathcal{D}} \left[ \left( Q_{\phi_i}(s, a) - y(r, s', d) \right)^2 \right],$$

The policy should, in each state, act to maximize the expected future return plus expected future entropy:

$$\begin{aligned} V^\pi(s) &= \mathbb{E}_{a \sim \pi} [Q^\pi(s, a)] + \alpha H(\pi(\cdot|s)) \\ &= \mathbb{E}_{a \sim \pi} [Q^\pi(s, a) - \alpha \log \pi(a|s)]. \end{aligned}$$

To get the policy loss, the final step is that we need to substitute  $Q^{\pi_{\{\theta\}}}$  with one of our function approximators

$$\max_{\theta} \mathbb{E}_{\substack{s \sim \mathcal{D} \\ \xi \sim \mathcal{N}}} \left[ \min_{j=1,2} Q_{\phi_j}(s, \tilde{a}_{\theta}(s, \xi)) - \alpha \log \pi_{\theta}(\tilde{a}_{\theta}(s, \xi)|s) \right],$$

# Proximal Policy Optimization (PPO)

PPO's combination of stability, sample efficiency, and compatibility with continuous action spaces has made it a popular choice for various reinforcement learning applications.

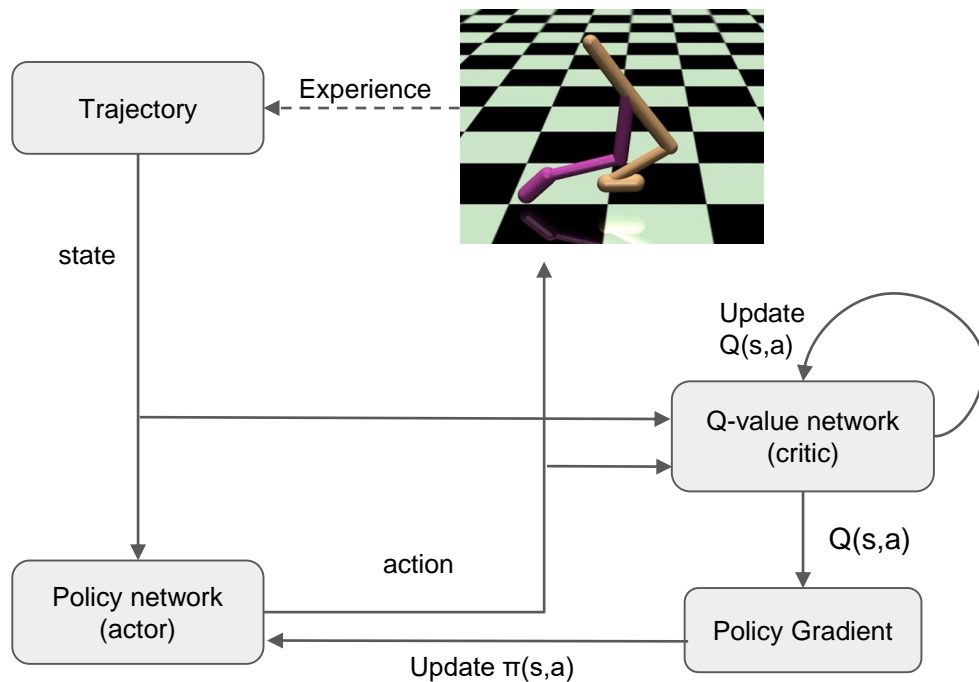
## Key Features:

- Works well with continuous action spaces, ideal for tasks that require fine-grained control.
- Policy updates are based on the most recent and relevant experiences as it is an on-policy algorithm.
- Objective function trades off between maximizing expected rewards and controlling the policy update magnitude.
- Incorporates a trust region approach to policy updates.
- Achieves effective policy learning with fewer interactions with the environment.

## Applications:


- PPO's ability to handle continuous action spaces and stability during training makes it suitable for complex control tasks in robotics.
- It has been used to achieve high-performance agents in challenging game environments, such as Atari games and complex strategy games.
- PPO is commonly used in simulated environments like the OpenAI Gym and MuJoCo.
- PPO has also been widely used as a benchmark algorithm in reinforcement learning research.

## Proximal Policy Optimization (PPO)





# Proximal Policy Optimization (PPO)



Ratio between the new policy and old policy

$$r_t(\theta) = \pi_{\theta}(a_t|s_t)/\pi_{\theta_k}(a_t|s_t)$$


Clipped PPO Objective Function

$$\mathcal{L}_{\theta_k}^{CLIP}(\theta) = \mathbb{E}_{\tau \sim \pi_k} \left[ \sum_{t=0}^T \left[ \min(r_t(\theta) \hat{A}_t^{\pi_k}, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t^{\pi_k}) \right] \right]$$


Policy Update

$$\theta_{k+1} = \arg \max_{\theta} \mathcal{L}_{\theta_k}^{CLIP}(\theta)$$

## Results




SAC Trained Walker2D Agent



PPO Trained Walker2D Agent

## Results



SAC Running Reward Plot  
(Training Time = 3476 seconds  
1,500,000 timesteps)



PPO Running Reward Plot  
(Training Time = 5787.92 seconds  
1,000,000 timesteps)

## Conclusion

- Both algorithms performed well on the environment, but PPO produced more stable results when compared to SAC based on the running rewards plot. Whereas the SAC experiment produced a much better result when compared based on the simulation generated.
- PPO being an on-policy algorithm, made it slower and less sample efficient compared to SAC, which is an off-policy algorithm.
- SAC, with its off-policy nature and soft value function updates, has a much better sample efficiency and faster learning on continuous control tasks like Mujoco Walker2d.
- PPO is generally regarded as a more stable algorithm and was easier to implement and tune compared to SAC, which has additional complexity due to the value function estimation.

# Thank You